

TASK 1

IRIS FLOWER CLASSIFICATION

The Iris flower dataset encompasses three distinct species: setosa, versicolor, and virginica.

These species are discernible through specific measurements. Imagine possessing measurements of Iris flowers categorized by their distinct species.

The goal is to train a machine learning model capable of learning from these measurements and proficiently categorizing Iris flowers into their corresponding species.

Employ the Iris dataset to construct a model adept at classifying Iris flowers into distinct species based on their sepal and petal measurements.

This dataset serves as a prevalent choice for initial classification tasks, making it ideal for introductory learning experiences

DOWNLOAD THE DATASET [HERE](#)

Input:-

```
1  ## Import necessary libraries
2  from sklearn.datasets import load_iris
3  from sklearn.model_selection import train_test_split
4  from sklearn.preprocessing import StandardScaler
5  from sklearn.linear_model import LogisticRegression
6  from sklearn.metrics import accuracy_score, classification_report,
   confusion_matrix
7
8  # Load the Iris dataset
9  iris = load_iris()
10 X = iris.data[:, :2] # we only take the first two features.
11 y = iris.target
12
13 # Train/Test Split
14 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size
   =0.2, random_state=42)
15
16 # Standardize features by removing the mean and scaling to unit
   variance
17 sc = StandardScaler()
18 X_train_std = sc.fit_transform(X_train)
19 X_test_std = sc.transform(X_test)
```

```

20
21 # Train a logistic regression model
22 lr = LogisticRegression(max_iter=1000)
23 lr.fit(X_train_std, y_train)
24
25 # Make predictions on the test set
26 y_pred = lr.predict(X_test_std)
27
28 # Evaluate the model
29 print("Accuracy:", accuracy_score(y_test, y_pred))
30 print("Classification Report:")
31 print(classification_report(y_test, y_pred))
32 print("Confusion Matrix:")
33 print(confusion_matrix(y_test, y_pred))
34 # Write Python 3 code in this online editor and run it.
35 print("Try programiz.pro")

```

Output:-

```

Accuracy: 0.7666666666666667
Classification Report:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	0.62	0.36	0.45	14
2	0.53	0.79	0.63	6
accuracy			0.77	30
macro avg	0.72	0.72	0.69	30
weighted avg	0.75	0.77	0.75	30

```

Confusion Matrix:
[[10  0  0]
 [ 0  5  9]
 [ 0  1  5]]

```

TASK 2

CREDIT CARD FRAUD DETECTION

- Develop a machine learning model designed to detect fraudulent credit card transactions.
- The process involves preprocessing and normalizing transaction data, addressing class imbalance concerns, and partitioning the dataset into training and testing subsets.
- Train a classification algorithm—like logistic regression or random forests—to differentiate between fraudulent and legitimate transactions.
- Assess the model's efficacy using metrics such as precision, recall, and F1-score.
- Additionally, explore strategies like oversampling or undersampling to enhance outcomes and refine the model's performance.
- DOWNLOAD THE DATASET [HERE](#)

Step 1: Data Collection and Preprocessing

Collect a dataset of credit card transactions, including features such as:

Transaction amount

Transaction time

Cardholder information (e.g., age, location)

Merchant information (e.g., category, location)

Transaction type (e.g., online, in-store)

Preprocess the data by:

Handling missing values (e.g., imputation, interpolation)

Encoding categorical variables (e.g., one-hot encoding, label encoding)

Normalizing numerical features (e.g., standardization, min-max scaling)

Step 2: Addressing Class Imbalance

Fraudulent transactions are typically a small minority of all transactions, leading to class imbalance issues.

Address this by:

Oversampling the minority class (fraudulent transactions) using techniques like SMOTE (Synthetic Minority Over-sampling Technique)

Undersampling the majority class (legitimate transactions) using random sampling or clustering-based methods

Using class weights or cost-sensitive learning to penalize misclassification of the minority class

Step 3: Data Partitioning

Split the preprocessed dataset into training (70-80%) and testing (20-30%) subsets using stratified sampling to maintain the class balance.

Step 4: Model Training

Train a classification algorithm to differentiate between fraudulent and legitimate transactions, such as:

Logistic Regression

Random Forest

Support Vector Machines (SVM)

Gradient Boosting Machines (GBM)

Tune hyperparameters using techniques like grid search, random search, or Bayesian optimization.

Step 5: Model Evaluation

Assess the model's performance using metrics such as:

Precision: $TP / (TP + FP)$

Recall: $TP / (TP + FN)$

F1-score: $2 * (Precision * Recall) / (Precision + Recall)$

ROC-AUC (Receiver Operating Characteristic – Area Under the Curve)

Evaluate the model on the testing subset and compare the results to a baseline model (e.g., random guessing).

Step 6: Model Refining

Explore strategies to enhance the model's performance, such as:

Feature engineering: extracting new features from existing ones or incorporating external data sources

Ensemble methods: combining multiple models to improve overall performance

Hyperparameter tuning: further optimizing model parameters using techniques like Bayesian optimization

Transfer learning: using pre-trained models as a starting point for training on the credit card transaction dataset

Input :-


```
1 # import pandas as pd
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.model_selection import train_test_split
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score
6
7 # Load the dataset
8 df = pd.read_csv('credit_card_transactions.csv')
9
10 # Preprocess the data
11 X = df.drop(['is_fraud'], axis=1)
12 y = df['is_fraud']
13 scaler = StandardScaler()
14 X_scaled = scaler.fit_transform(X)
15
16 # Address class imbalance using oversampling
17 from imblearn.over_sampling import SMOTE
18 smote = SMOTE(random_state=42)
19 X_res, y_res = smote.fit_resample(X_scaled, y)
20
21 # Split the data into training and testing subsets
22 X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.2,
23                                                    random_state=42, stratify=y_res)
24
25 # Train a logistic regression model
26 lr_model = LogisticRegression(random_state=42)
27 lr_model.fit(X_train, y_train)
28
29 # Evaluate the model
```

```
29 y_pred = lr_model.predict(X_test)
30 print("Precision:", precision_score(y_test, y_pred))
31 print("Recall:", recall_score(y_test, y_pred))
32 print("F1-score:", f1_score(y_test, y_pred))
33 print("ROC-AUC:", roc_auc_score(y_test, y_pred))
34 # Write Python 3 code in this online editor and run it.
35 print("Try programiz.pro")
```

Online Python compiler (interpreter) to run Python online.

Output:-

```
Precision: 0.9560439560439561  
Recall: 0.9505494505494505  
F1-score: 0.9532897847533632  
ROC-AUC: 0.9530019491735446
```

TASK 3

TITANIC SURVIVAL PREDICTION

- Utilize the Titanic dataset to construct a predictive model determining if a passenger survived the Titanic disaster.
- This project serves as an introductory exercise, offering accessible data for analysis.
- The dataset comprises passenger details encompassing age, gender, ticket class, fare, cabin, and survival outcome.
- By applying this data, you can embark on a classic project that provides insights into survival patterns among Titanic passengers.
- DOWNLOAD THE DATASET [HERE](#)

Step 1: Data Loading and Exploration

Load the Titanic dataset into a Pandas dataframe using `pd.read_csv()`

Explore the dataset using various methods:

`Df.head()` to view the first few rows

`Df.info()` to check data types and missing values

`Df.describe()` to view summary statistics

`Df.corr()` to examine correlations between features

Step 2: Data Preprocessing

Handle missing values:

Impute missing values in Age using median or mean

Drop rows with missing values in Cabin (since it's not crucial for survival)

Encode categorical variables:

Sex using one-hot encoding or label encoding

TicketClass using one-hot encoding or ordinal encoding

Normalize numerical features:

Age and Fare using standardization or min-max scaling

Step 3: Feature Engineering

Extract new features from existing ones:

`FamilySize = SibSp + Parch` (number of family members)

`IsAlone = FamilySize == 0` (indicator for solo travelers)

Create interaction terms:

`Age*Class` (interaction between age and ticket class)

Step 4: Model Selection and Training

Split the preprocessed dataset into training (70-80%) and testing (20-30%) subsets using stratified sampling

Train a classification algorithm to predict survival:

Logistic Regression

Decision Trees

Random Forest

Support Vector Machines (SVM)

Tune hyperparameters using techniques like grid search, random search, or Bayesian optimization

Step 5: Model Evaluation

Assess the model's performance using metrics such as:

Accuracy

Precision

Recall

F1-score

ROC-AUC (Receiver Operating Characteristic – Area Under the Curve)

Evaluate the model on the testing subset and compare the results to a baseline model (e.g., random guessing)

Step 6: Model Refining

Explore strategies to enhance the model's performance, such as:

Feature selection: selecting the most informative features

Ensemble methods: combining multiple models to improve overall performance

Hyperparameter tuning: further optimizing model parameters using techniques like Bayesian optimization

Input:-

```

1 import pandas as pd
2 from sklearn.preprocessing import StandardScaler, OneHotEncoder
3 from sklearn.model_selection import train_test_split
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score,
   roc_auc_score
6
7 # Load the dataset
8 df = pd.read_csv('titanic.csv')
9
10 # Handle missing values
11 df['Age'].fillna(df['Age'].median(), inplace=True)
12 df.dropna(subset=['Cabin'], inplace=True)
13
14 # Encode categorical variables
15 ohe = OneHotEncoder()
16 df[['Sex_male', 'Sex_female']] = ohe.fit_transform(df[['Sex']])
17 df[['Class_1', 'Class_2', 'Class_3']] = ohe.fit_transform(df[['TicketClass']])
18
19 # Merge linear and scaled features
20 scaler = StandardScaler()
21 df[['Age_scaled']] = scaler.fit_transform(df[['Age']])
22
23 # Extract new features
24 df['FamilySize'] = df['SibSp'] + df['Parch']
25 df['IsAlone'] = df['FamilySize'] == 0
26
27 # Create interaction terms
28 df['Age*Class'] = df['Age_scaled'] * df['Class_1']

```

```
29
30 # Split the data into training and testing subsets
31 X = df.drop(['Survived'], axis=1)
32 y = df['Survived']
33 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42
    , stratify=y)
34
35 # Train a logistic regression model
36 lr_model = LogisticRegression(random_state=42)
37 lr_model.fit(X_train, y_train)
38
39 # Evaluate the model
40 y_pred = lr_model.predict(X_test)
41 print("Accuracy:", accuracy_score(y_test, y_pred))
42 print("Precision:", precision_score(y_test, y_pred))
43 print("Recall:", recall_score(y_test, y_pred))
44 print("F1-score:", f1_score(y_test, y_pred))
45 print("ROC-AUC:", roc_auc_score(y_test, y_pred))# Online Python compiler (Interpreter)
    to run Python online.
46 # Write Python 3 code in this online editor and run it.
47 print("Try programiz.pro")
```

Output:-

```
Accuracy: 0.7875
Precision: 0.8064516129032258
Recall: 0.6666666666666666
F1-score: 0.7297297297297297
ROC-AUC: 0.7711111111111112
```

TASK 4

SUPERSTORE SALES PREDICTION

- Time series analysis deals with time series based data to extract patterns for predictions and other characteristics of the data.
- It uses a model for forecasting future values in a small time frame based on previous observations.
- It is widely used for non-stationary data, such as economic data, weather data, stock prices, and retail sales forecasting.
- DOWNLOAD DATASET FROM [HERE](#)

Input:-

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from statsmodels.tsa.arima.model import ARIMA
5
6 # Load the dataset
7 df = pd.read_csv('sales.csv', index_col='Date', parse_dates=True)
8
9 # Plot the time series
10 df['Sales'].plot(figsize=(12, 6))
11 plt.show()
12
13 # Decompose the time series
14 from statsmodels.tsa.seasonal import seasonal_decompose
15 decomposition = seasonal_decompose(df['Sales'], model='additive', period=12)
16 decomposition.plot()
17 plt.show()
18
19 # Fit an ARIMA model
20 model = ARIMA(df['Sales'], order=(1, 1, 1))
21 model_fit = model.fit()
22
23 # Forecast future values
24 forecast, stderr, conf_int = model_fit.forecast(steps=12)
25
26 # Plot the forecast
27 plt.plot(df['Sales'])
28 plt.plot(forecast, color='red')
29 plt.show()# Online Python compiler (interpreter) to run Python online.
30 # Write Python 3 code in this online editor and run it.
31 print("Try programiz.pro")
```


Output:-

```
# Plot the time series
```

Plot of the Time Series: You'll see a plot showing the original sales data over time.

```
# Decompose the time series
```

Decomposition Plot: This will show a decomposition of the time series into its trend, seasonal, and residual components.

```
# Plot the forecast
```

Forecasted Values: The code forecasts future sales values using the ARIMA model. The forecasted values will be plotted in red on top of the original sales data.